

PROJECT TITLE: Automating Data Quality Monitoring in Cloud Warehouses with AI

College Name: HKBK College of Engineering

Group Members:

- **Name:** Kavya R Naik
CAN ID Number: CAN_33740336
 - **Name:** Md Modassir Ali
CAN ID Number: CAN_33506401
 - **Name:** Dhakshana Moorthy
CAN ID Number: CAN_33836237
-

1. Abstract

This project introduces an AI-powered solution to automate data quality monitoring in cloud data warehouses. Traditional manual monitoring methods are time-intensive and prone to human error, failing to scale with the increasing complexity of modern data environments. By automating data quality monitoring in cloud data warehouses using artificial intelligence (AI), this overcomes the drawbacks of manual techniques that are laborious, prone to mistakes, and unable to keep up with the increasing complexity of contemporary data environments. By using cloud platforms like AWS, Google Cloud, or IBM Cloud, together with tools like Python and Scikit-learn, the system finds and fixes problems with data quality such as anomalies, inconsistencies, and schema deviations. The solution guarantees data dependability, offers actionable insights through real-time dashboards and alerts, reduces manual involvement, and effortlessly integrates into organizational data workflows, all of which contribute to increased confidence in data-driven decisions.

2. Problem Definition

Data Problem

Cloud data warehouses store vast amounts of data from various sources. Ensuring the quality of this data is critical for downstream analytics and business intelligence. However, existing methods for monitoring and ensuring data quality are manual, time-consuming, and error-prone. This leads to:

- **Delayed detection of issues:** Late identification of missing data, duplicates, or anomalies.
- **Inefficiency:** Increased reliance on manual interventions for validating data quality.
- **Reduced trust in analytics:** Poor-quality data negatively impacts decision-making.

Key challenges include:

- Detecting real-time anomalies in data volume, schema, and content.
- Automating validation processes to reduce manual overhead.
- Providing actionable insights for addressing quality issues.

Key Question:

- How can organizations effectively monitor and maintain data quality in cloud data warehouses?
- Can AI automate the identification and resolution of common data quality issues such as missing values, anomalies, or duplicates?

Target User:

- **Data Engineers:**
 - **Responsibilities:** Maintain data pipelines and ensure data integrity.
 - **Needs:** Tools to automate anomaly detection and streamline validations.
- **Data Analysts:**
 - **Responsibilities:** Extract insights from clean, structured data.
 - **Needs:** Reliable data quality metrics to trust analytical outcomes.
- **Business Intelligence Teams:**
 - **Responsibilities:** Generate reports and dashboards for strategic decision-making.
 - **Needs:** High-quality data inputs to ensure accurate reporting.

Goal:

- Automate data quality checks to reduce manual interventions.
 - Build a scalable solution to monitor data patterns, schema changes, and anomalies in real-time.
 - Provide actionable insights via interactive dashboards and alert systems to improve operational efficiency.
-

3. Requirements

Functional Requirements:

- **Data Ingestion and Processing:**
 - Load data from cloud data warehouses (e.g., AWS Redshift, Google BigQuery).
 - Clean and preprocess data to handle missing values and duplicates.
- **Anomaly Detection:**
 - Implement machine learning models for detecting data anomalies.
 - Identify issues in data volumes, patterns, and schema.
- **Reporting and Alerts:**
 - Generate real-time alerts for data quality issues.
 - Provide detailed reports on data quality trends and anomalies.
- **Integration:**
 - Expose the solution as an API for integration with existing workflows.

Non-Functional Requirements:

Non-Functional Requirements

- **Scalability:**
 - Ensure the system handles high transaction volumes from enterprise-grade data warehouses.
- **Security:**
 - Employ robust encryption for data in transit and at rest to maintain confidentiality.
- **Performance:**
 - Detect anomalies with low latency to support real-time workflows.
- **Reliability:**
 - Provide consistent uptime and error resilience for continuous monitoring.

4. Tools and Platforms

- **Data Preprocessing Tools:**
 - Python, Pandas for cleaning and transforming data.
- **Model Development and Training:**
 - Scikit-learn for implementing anomaly detection algorithms.
 - TensorFlow or PyTorch for building advanced models if needed.

- **Visualization and Reporting:**
 - Matplotlib, Seaborn for generating visualizations.
 - Power BI, Tableau, or Looker for building interactive dashboards.
 - **Cloud Platforms:**
 - **AWS:** S3 for storage, SageMaker for model training, Lambda for serverless API integration.
 - **Google Cloud:** BigQuery for querying, Vertex AI for training and deploying models.
 - **IBM Cloud:** Object Storage and Watson Studio for data storage, preprocessing, and model development.
 - **Monitoring and Logging:**
 - Tools like Grafana and CloudWatch for performance monitoring and anomaly alerts
-

5. Implementation Plan

Step 1: Data Preparation:

- Load sample datasets from the chosen cloud data warehouse.
- Perform initial data profiling to assess data quality metrics.
- Clean and preprocess the data to handle missing values, outliers, and duplicates.

Step 2: Model Development:

- Train an anomaly detection model using Scikit-learn (e.g., Isolation Forest) or TensorFlow.
- Evaluate model performance using metrics such as precision, recall, and F1-score.
- Optimize the model by tuning hyperparameters and testing on different datasets.

Step 3: Deployment:

- Deploy the AI model as a REST API using IBM Watson Machine Learning.
- Deploy the trained model to a cloud platform (e.g., AWS SageMaker, Google Vertex AI, or IBM Watson Studio).
- Expose the model as an API for real-time anomaly detection.
- Test the API with sample data to ensure proper integration.
- Integrate the API with the cloud warehouse for continuous monitoring.

Step 4: Reporting and Visualization:

- Use Matplotlib or Seaborn to visualize trends and flagged anomalies.
 - Create interactive dashboards using Power BI or Tableau to monitor data quality in real-time.
 - Generate automated reports with actionable insights.
-

6. Expected Outcomes

- A fully functional AI-driven system capable of automating data quality monitoring.
- An API for seamless integration with existing workflows.
- Reduced operational costs associated with data quality management.
- Interactive dashboards and reports for real-time insights into data quality trends and issues.
- Enhanced efficiency and reliability in managing data quality, reducing manual efforts, and improving trust in analytics.